

MedFusion: A Unified AI Framework for Patient Interaction and Clinical Intelligence

R Hemanth Kumar¹, Sachin R², Vishnu Prasad B³, Himanshu Yadav⁴, Sajitha Krishnan^{5*}

Department of Computer Science and Engineering

Amrita School of Computing, Bengaluru

Amrita Vishwa Vidyapeetham, India

¹bl.en.u4aie22138@bl.students.amrita.edu, ²bl.en.u4aie22142@bl.students.amrita.edu,

³bl.en.u4aie22164@bl.students.amrita.edu, ⁴bl.en.u4aie22167@bl.students.amrita.edu,

⁵k_sajitha@blr.amrita.edu

Corresponding author: Sajitha Krishnan, k_sajitha@blr.amrita.edu

Abstract—At present, healthcare professionals experience inefficiency due to the disconnect between clinical documentation and diagnostic imaging analyses. "MedFusion" is an innovative AI-based platform that combines diagnostic imaging with verbal patient/physician interactions using a single source of capture. The MedFusion platform employs a sequential multi-modal pipeline. The conversation captures audio from patient interviews through the use of OpenAI's WHISPER, then processes the data through the MTS-Dialog Dataset to extract factual content. In parallel, the image analysis captures images with a Vision Transformer Encoder and a DistilGPT2 Decoder. Objective radiologic findings for the patient are extracted from the images and enhanced with knowledge gained from the Radiology Objects in Context (ROCO) dataset. The major difference between MedFusion and current multimodal healthcare AI frameworks using end-to-end generation or loosely coupled pipelines is that MedFusion proposes a deterministic fusion engine, which guarantees the clinical safety by producing no hallucinations during the integration of structured medical images and unstructured conversation data. MedFusion characterizes the capture and extracts/constructs the reports based on the software systems used to analyze and process the images. Experimental results revealed that, measured using the ROUGE-L score, the ViT-GPT2 architecture obtained a score of 0.3, and that the training loss for the architecture converged to below 0.1 over 20 epochs. Additionally, the report assembly (fusing) by the MedFusion Engine occurred in less than 0.01 seconds. MedFusion combines capabilities that allow a tool to simultaneously "see" and "hear" from the same patient data, rather than creating new information, and uses the software developed using PyTorch and Python to drive the MedFusion platform to help alleviate the administrative burden placed on healthcare practitioners and increase diagnostic accuracy.

Index Terms—Multimodal AI, Medical Imaging, Clinical Decision Support, Speech Recognition, Healthcare AI

I. INTRODUCTION

In this modern world, AI technology has greatly contributed to the development of the healthcare sector, facilitating automated diagnosis, imaging analysis with steady ongoing research on these domains [16-20].

The process of treatment within the medical industry is multimodal and requires doctors to analyze both structured and unstructured information at the same time. The current AI tools mainly focus on analyzing one modality at a time without considering the interactions between narrative accounts and diagnostic images. This problem leads to increased inefficiencies, cognitive loads for practitioners, and chances of misdiagnosis within the industry.

A literature survey reveals that although many algorithms have proven successful in accomplishing single modality applications like radiology image captioning and clinical text summarization, there are no existing methods that can efficiently combine multiple clinical modalities in a coherent real-time decision support system. The current techniques involve either loosely coupled pipeline-based approaches or end-to-end generation-based frameworks, both of which possess certain shortcomings such as being non-interpretable, clinically unreliable, and prone to hallucinations. This poses an important issue in the field of medicine because accurate correlation of a patient's history and radiology results is imperative for making correct diagnostic decisions.

MedFusion is an innovative AI platform that attempts to solve the problem of disparate information sources by allowing doctors and patients to have a single conversation within the medical image analysis process. MedFusion integrates many aspects of medical technology into one single platform for both narrative data (from doctors and patients' spoken conversations) and visual diagnostic evidence (from images). A key feature of MedFusion is that it uses a sophisticated multi-modal, sequential pipeline to do this. First, it takes audio input from conversations and converts that into text using speech-to-text technology, in order to extract factual information related to the conversation with the help of a localized LLM. The second part of this pipeline, occurring simultaneously, uses the medical images to provide objective findings

about the patient's condition. By creating a single integrated tool that can process images ("see"), audio ("hear"), and combine all the information ("assemble") in a deterministic way, MedFusion is able to provide new and unique approaches to using clinical informatics. This should allow for improved quality of patient care by reducing the chance for error due to "hallucinations" that occur when using only generative models. In practice, when dealing with real-world settings, it is inevitable for doctors to make interpretations based on the information available from different channels; thus, it can be difficult for them to notice some crucial connections between the signs observed and the evidence collected. Hence, it has become necessary to develop an efficient system that could help integrate different inputs into a useful output.

II. RELATED WORKS

[1] A Transformer based architecture combining a memory augmented Sparse Attention block and a Medical Concepts Generation Network (MCGN) is proposed. This architecture is tested on the extensive MIMIC-CXR dataset with CLIP-B/16 serving as the visual feature extractor. The authors note that there is still a problem in effectively incorporating even more intricate and nuanced medical concepts to improve the accuracy of generated reports, especially in cases that can be challenging or unclear.

[2] An Inclusive Task Aware architecture is introduced to improve the quality of reports by utilizing a task distillation module, a categorization token, and a multi-head Transformer based report generator. The model is evaluated on the IU X-ray and MIMIC-CXR datasets. Improving the model's ability to generate accurate descriptions for rare anomalies and ensuring all relevant anatomical features are reliably and consistently reported in each generated summary both of which are critical for comprehensive diagnostics remain major challenges.

[3] A self-guided framework (SGF) employs a Transformer Report Generator and a Knowledge Distiller to improve report quality in an unsupervised. The model used is ResNet-101 for visual extraction and Sentence-BERT for text embedding. This framework is a model which allows it to generate its own outputs without direct supervision while using an unsupervised. Improving this self-guidance mechanism is a strong research gap.

[4] MedBLIP demonstrates strong zeroshot classification and visual question answering abilities without the need for task-specific training. This makes possible to seamlessly combine clinical text with 3D images. This approach is highly adaptable and efficient. The authors suggest that future studies broaden this bootstrapping approach to incorporate a wider range of 3D medical imaging modalities and structured Electronic Health Record data in order to

create more comprehensive and powerful diagnostic support systems.

[5] This study proposes XPRONET a network driven by a shared cross-modal prototype matrix, to enhance the alignment of visual data from images with textual elements in reports. Using the 'IU-Xray' and MIMIC-CXR datasets, the model extracts image features using a Transformer encoder-decoder backbone and ResNet-101. Primary innovation is the use of an improved multi-label contrastive loss which enhances the prototype's learning process and leads to a more successful alignment of the visual and textual data modalities. One of the major unresolved issues in radiology is how to enhance this cross modal prototype learning to better capture the complex, deep, and often subtle relationships between visual features and the textual descriptions that go along with them.

[6] The model is guided by the Segment Anything Model, and MSMedCap generates detailed captions for medical images. Using a dual-encoder architecture that blends general features from CLIP with specific, fine-grained details from SAM, it applies this to datasets like ROCO and MedICaT. The text is generated using an OPT Large Language Model. The overall accuracy and level of detail of the final analysis are enhanced by the model pretraining ability to use mixed semantic learning to gather both general and highly specialized medical image features.

[7] Research in this area develops Flamingo-CXR by improving the Flamingo vision-language foundation model. The study makes use of the private Indian dataset IND1 as well as the public MIMIC-CXR dataset. The primary finding of this study is that physicians often prefer AI generated report, underscoring the significance of clinician AI collaboration in improving the overall general quality, accuracy and efficacy of reports in clinical procedures..

[8] The robust system design includes specialized classification modules for various diagnostic tasks. In order to ensure that the generated reports are clinically reliable and useful the study emphasizes the critical need for evaluation metrics that assess both language fluency and diagnostic accuracy. It accomplishes by utilizing knowledge from datasets like CheXpert and MIMIC-CXR. The author's primary research gap is figuring out how best to integrate these various specialized modules.

[9] The ROCov2 paper presents an updated multimodal dataset that includes over 81,000 radiological images from PubMed publications with their original captions and their corresponding CUI IDs. For CT, MRI and X-ray, this dataset provides accurate annotations on imaging types, anatomical regions, and X-ray directional information. Pretraining more powerful models could be done for radiology report production using large, rich, and well-annotated

datasets such as ROCov2.

[10] The suggested Region-Guided Radiology Report Generation (RGRG) model segments using a Faster R-CNN with a ResNet-50 using a refined GPT-2 Medium. Using the Chest ImaGenome dataset, which is extracted from MIMIC-CXR. Future work of this paper suggests to improve the segmentation module in order to accurately identify all clinically relevant regions, especially in complex situations.

[11] The MedVAG (Medical Vision Attention Generation) framework, which combines a GPT-2 model for language generation and a Vision Transformer (ViT) for visual feature extraction, is trained on the IU X-Ray and COV-CTR datasets. The model is enhanced with graph-based feature fusion, cross-attention, and memory-guided attention.

[12] A prototype system uses symptoms to predict diseases. Multiple ML classifiers is trained using patient encounter reports. Their approach combines traditional symptom analysis for diagnosis with Vader based sentiment analysis of drug reviews. Challenges remains on how to effectively merge these multiple data sources—symptoms for reliable recommendation without the end result being altered by inconsistent user ratings.

[13] Support Vector Classifier and Random Forest (RF) models are trained on dataset from Kaggle containing 4,920 samples to predict diseases. A unique feature of this is that it provides a more complete support system for patients. The research gap identified by the authors is a critical one for real-world application: how to ensure the reliability and safety of these recommendations.

[14] Predicts both general diseases based on symptoms using models such as Naive Bayes, Random Forest. Its integrated module, suggests medical specialists to the user based on the anticipated ailments and their particular data, is a distinctive feature.

[15] Using a Yelp dataset from the "Health & Medical" category, the RR&R-CNN model which combines a CNN with a factorization machine is proposed to generate medical suggestions. It vectorizes text using BERT and updates user ratings based on sentiment analysis using TextBlob. The open problem it addresses is basic barrier in user generated data for the emotional bias and inherent inconsistencies in user provided data.

III. OUR CONTRIBUTIONS


The proposed framework uses a deterministic approach to multi-modal fusion as compared to the traditional generative and black-box approaches to multi-modal fusion. In contrast to previous work where either modality independence or probability fusion was applied, the MedFusion method guarantees transparency in the fusion of clinical modalities with reduced risk of hallucination.

Our contribution include :

- **Unified Multimodal Framework:** A single interface to process both structured visual data (X-ray images) and unstructured narrative data (patient-physician audio) to close the digital gap between radiology and clinical history by combining OpenAI's Whisper model for speech transcription with a Vision Transformer (ViT) in conjunction with DistilGPT-2 for image captioning.
- **Deterministic Rule-Based Fusion:** While traditional generative fusion systems tend to incorporate hallucinations in their report generation process, the proposed approach ensures that no such issues arise by fusing only clinically validated output from each modality within specified templates.
- **Less Hallucination Fact Extraction:** Utilizing LLMs specifically for entity extraction (e.g., symptoms, vitals) rather than open-ended generation. The system retains relevant terminologies.

IV. DATASET

A. MTS dialog dataset



section_text string : length	dialogue string : length
372*533 9.9%	41*924 80.9%
Symptoms: no fever, no chills, no cough, no congestion, no nausea, no vomiting, no chest pain, no chest pressure. Diagnosis: hypertension, osteoarthritis, osteoporosis, hypothyroidism, allergic rhinitis, kidney stones History of Patient: 76-year-old white female, presents to the clinic today originally for hypertension and a med check, followed by Dr. Kumar, issues stable Plan of Action: N/A	Doctor: what brings you back into the clinic today, miss? Patient: I came in for a refill of my blood pressure medicine. Doctor: It looks like Doctor Kumar followed up with you last time regarding your hypertension, osteoarthritis, osteoporosis, hypothyroidism, allergic rhinitis and kidney stones. Have you noticed any changes or do you have any concerns regarding these issues? Patient: No. Doctor: Have you had any fever or chills, cough, congestion, nausea, vomiting, chest pain, chest pressure? Patient: No. Doctor: Great. Also, for our records, how old are you and what race do you identify yourself as? Patient: I am seventy six years old and identify as a white female.

Fig. 1. Sample entries from the MTS Dialog Dataset showing the structure of doctor-patient conversations.

The MTS Dialog dataset Fig. 1. is a collection of structured conversations involving doctors and patients for the purpose of creating a treatment plan. Extensive exploratory data analysis took place on each of the three separate data sets—training, validation, and test sets—of dialogue data, which is composed of multiple columns, such as ID, section_text, diagnosis, and treatment_plan.

In order to balance the class imbalance found in the MTS Dialog Dataset, the following measures have been taken for pre-processing and augmentation. The minority classes have been upsampled by means of various data augmentation methods such as paraphrasing and synonym replacement. In addition, the class weighing approach has been used in the process of training the models where the misclassified classes get more weightage during the training period.

The exploratory data analysis for the dataset consisted of performing several analyses on the three

files, such as examining the overall structure of the dataset, understanding the characteristics of each data type (e.g., text, number), examining the number of null values in each dataset, and ensuring that all datasets are complete and consistent. The examination of these characteristics demonstrates that there is a large variety of medical topics being addressed through this dataset, and that there is a great deal of variation in the sentence length and tone of the conversation between the doctor and the patient.




Fig. 2. MTS section header distribution showing the frequency of clinical sections

Fig. 2 displays the frequency of top 9 medical vocabulary occurrences in the dataset that determine the length of physical and symptoms-based conversations, identify patterns that exist between symptoms and diagnoses, etc. This dataset contains a wealth of medically contextual conversation data, however, the dataset is highly imbalanced in the amount of various diagnosis types being represented in the dataset, as well as being unevenly distributed across the disease category.

The knowledge obtained from exploring the dataset is foundational for developing Natural Language Processing-based or Dialogue Systems for Automated Medical Consultations, as the data requires normalization and class augmentation in order to account for the variances and class imbalance.

B. ROCO dataset

The ROCO (Radiology Objects in Context) Dataset Fig. 3. is intended to connect visual and textual medical data by combining radiological images and their respective text descriptions (i.e. captioning). The exploratory data analysis starts with a clean bagged record of the data and an overview of all of the image/caption combinations, types of data within them, and missing values.



Fig. 3. Sample X-ray images from the ROCO dataset

The analysis reviews structural characteristics of captions, most frequently used words, and frequency of occurrence of terms specific to radiology. This analysis uses tools such as histograms, word clouds, and bar plots to summarize findings on the most frequently included organ systems, modalities (e.g. X-ray, MRI, CT), and diseases/conditions noted in captions.




Fig. 4. Word cloud visualization of the most frequent anatomical and modality terms found in radiology captions.

In addition to the above, the analysis describes the length distributions of captions and diversity of the medically specific terms. Fig. 4 presents a word cloud visualization of the caption content, visually highlighting the high frequency of anatomical terms (e.g., "lung," "spine") and modality indicators (e.g., "x-ray," "computed tomography") [9].




Fig. 5. Distribution of the top 15 Concept Unique Identifiers (CUIs)

Fig. 5. shows the distribution of CUIs from the Unified Medical Language System (UMLS) which

maps a specific medical concepts to a unique identifier. The distribution of these CUIs reveals a broad range of clinical areas, including Cardiology and Gastroenterology.

C. Audio Recording Dataset (Kaggle)

"Audio Recording Whisper Dataset" from Kaggle is used to guarantee the acoustic resilience of the MedFusion platform. The lack of comprehensive multimodal medical datasets with ground truth summaries, digital imaging, and raw audio is addressed by this addition. The raw audio waveforms needed to train and evaluate Automatic Speech Recognition (ASR) systems are absent from the MTS-Dialogue dataset, despite the fact that it offers precise text transcriptions.

As a result, the Audio Recording Whisper Dataset was chosen as the main input for assessing the platform's Speech-to-Text (STT) capabilities. This dataset functions as a specialized training set with the goal of establishing accurate recognition of medical terminology, a variety of accents, and different audio characteristics that are present in real-world user interactions. It includes an MP3 collection of 273 unique doctor-patient consultations and all the audio files are converted to 16kHz WAV format. This ensures perfect compatibility to the whisper model. The pipeline ensures asynchronous processing of modalities during feature extraction, followed by synchronized integration at the fusion stage, enabling both efficiency and consistency in report generation.

V. SYSTEM ARCHITECTURE

The proposed system operates on a multimodal pipeline, which is made up of three modules, namely, speech processing, image analysis, and deterministic fusion.




Fig. 6. System Architecture

Fig. 6 Shows that all three modules have been developed as independent components, where each one has an input and output specification. The two types of data, namely, audio and images, are processed independently during the feature extraction process and then fused together based on a set of rules. The

suggested approach has been implemented in Python using PyTorch library. This has two primary purposes:

(1) Resource optimization: Because the components of MedFusion are processed in a sequenced manner, the MedFusion platform can operate with ease on standard consumer grade hardware, including local GPUs, so the only limitations on processing speed should be the amount of Video RAM (VRAM) on the hardware.

(2) Clinical Safety: By assigning the "reasoning" from the clinical documents (the extraction and captioning of the information) to a compartmentalized task and allowing the final report generation to be completed using a logical inference engine, the potential exists that an AI-incorrect diagnosis created using reasoning will no longer be able to occur.

The system is composed of three primary functional blocks:

A. The Speech Processing Block

The Speech Processing Block transforms unstructured conversations between doctors and patients into clinical data. The input audio file can be either in .mp3 or .wav formats and is initially normalized and converted to 16kHz with the use of pydub package. Then, the output waveform is submitted to the Whisper ASR for generating a transcript of words spoken. The next step involves the processing of this transcript through a minimalistic language model for extracting medical facts in terms of NER (Chief Complaint, HPI, Symptoms, and Vital Signs). It is done by adopting a low-temperature decoding method ($T=0.1$).

B. The Image Analysis Block

The Image Analysis Block analyzes radiological images for producing clinically descriptive findings. The input images undergo preprocessing such as resizing and normalization and are split into fixed-sized patches. They are further fed into the ViT encoder, which extracts the global context between different parts of the images. These feature representations are then used by the DistilGPT-2 decoder to produce descriptive textual findings about the radiology images.

The Vision Transformer (ViT) architecture was chosen as the visual encoder since it was able to take advantage of global relationships among the pixels throughout the image, which is highly beneficial in medical applications when distant locations may correlate to certain conditions. The DistilGPT-2 decoder was selected due to the lightweight nature of this language model and efficient inference, allowing for deployment on limited resources.

C. The Deterministic Fusion Block

The Deterministic Fusion Block fuses the results extracted by the two modules into a final clinical report. As opposed to other generative fusion modules,

the deterministic approach involves rule-based fusion pipelines. The extracted text is firstly preprocessed through normalization of strings. Normalization involves lower casing, whitespace normalization, and punctuations. Regex techniques are further used to check whether the extracted entities are correct and follow certain formats. A predetermined template for a clinical report is used, with placeholders being replaced with data extracted by each of the modules.

VI. MEDFUSION: DATA FLOW AND PIPELINE LOGIC

The proposed system processes data through a distinct, transparent pipeline.




Fig. 7. Data-Flow Diagram

A. Front-end Data Flow Pipeline

The pipeline begins by accepting the raw input files (audio and images) Fig. 7. The system employs the pydub library to normalize incoming audio files which supports any format. The audio is then converted to a standardized 16kHz WAV format to ensure compatibility with the Whisper model.

The Speech Module (Whisper) takes the raw waveform of a patient's speech and converts it into text by using the rules of heuristic logic to create turns that are labeled as coming from either the Doctor or Patient.

The Image Module (Vision Transformer) extracts from it visual indications. The Image Module is not impacted or influenced by the Speech Module's output due to this parallel execution.

B. Rule-Based Synthesis

The proposed system deterministic assembler. It checks if the inputs are valid non-empty strings. Regex formatting is applied to ensure the structure of data that is being presented to the user. It is done systematically by inserting the text into specific placeholders: "HISTORY_SECTION" that contains the conversation summary and "IMAGING_SECTION" that contains the insights generated by the model. The model generates a dynamic assessment sentence by structuring the findings and the symptoms that is easy to read and interpret. This design guarantees that the final report is strictly a sum of its parts, ensuring zero hallucination.

VII. EVALUATION METRICS

In order to provide a thorough analysis of the suggested approach, two types of validations were utilized: component level and system level validations. The evaluation procedure will consist of not only the quantitative analysis of the speech and image components according to NLP and vision benchmarks but also qualitative and Case-based validations for the multimodal pipeline. The data sets were split into three partitions with an 80/10/10 split.

A. Speech Processing Module

Performance testing for the Speech Processing Module involved the use of a validation set that made up 10% of the MTS Dialog dataset. The generated summary outputs were compared to the summaries used as reference using the ROUGE measures (ROUGE-1, ROUGE-2, ROUGE-L, and ROUGE-SUM). The measures capture aspects of both lexical matching and semantic relevance within the context of the clinical information extraction task.

TABLE I
MODEL PERFORMANCE METRICS

Model	ROUGE-1	ROUGE-2	ROUGE-L	ROUGE-SUM
GPT2	0.68	0.55	0.62	0.62
GPT-NEO	0.63	0.55	0.62	0.63
DistilGPT-2	0.61	0.54	0.60	0.61

According to ROUGE summarization metrics from Table I, the model performed well in terms of summarization abilities. Specifically, ROUGE scores of GPT2 respectively, which indicates that generated summaries have a significant amount of content overlap with the original reference texts. The results show that the model is capable of capturing unigram-level information as well as bigram-level information when generating fluent and coherent summaries for longer text strings.

B. Image Processing Module

The Image Analysis Module was tested using metrics for evaluating generated descriptions of images, such as BLEU-1, BLEU-4, ROUGE-L, and METEOR. The results of these tests indicate the accuracy of the produced captions compared to ground truth radiology reports. This testing was performed based on a certain number of test cases selected from the ROCO dataset.

TABLE II
MODEL PERFORMANCE METRICS

Model	BLEU-1	BLEU-4	ROUGE-L	Meteor
ResNet50 LSTM	0.3	0.1	0.3	0.3
ResNet50 GPT-2	0.1	0.2	0.1	0.1
ResNet-101 T5	0.3	0.2	0.5	0.4
ViT (d)GPT-2	0.2	0.1	0.3	0.2
SWIN GPT-2	0.2	0.1	0.3	0.2

In the case of transformer models from Table II, which have produced lower values for BLEU and ROUGE, this could be explained by the fact that transformers produce paraphrases of the target descriptions, as opposed to exact matches with the reference text. Traditional metrics tend to place more emphasis on n-grams, while transformers are capable of producing semantically rich descriptions that cannot be accurately captured using the standard metrics.

Finally, the classic ResNet-50 + LSTM architecture has been established as a sound benchmark with lower-than-expected BLEU-4 scores (0.1) when generating captions and having greater susceptibility to confusion than modern neural network models with robust attention mechanisms such as Transformers. However, the results clearly show that while CNN-based encoders can yield excellent results based upon numerical evaluation metrics; the use of Transformer architecture allows for a much more favorable qualitative evaluation, which emphasizes the importance of assessing language generation performance using both automated numerical evaluations and qualitative human evaluations.

In addition to evaluating each individual module, the overall MedFusion system was tested using the method of Case-based validation. Each test case included an audio input (conversation between doctor and patient) along with its radiological counterpart. The criteria for success in testing included the accuracy of clinical entity extraction, accurate description of images, as well as correct integration of both types of information into a clinical report. In the absence of publicly accessible multimodal data sets, a qualitative evaluation of clinical reports was conducted.

VIII. DEMONSTRATION

The Rule-Based Fusion Engine synthesized the extracted text with the image findings.




Fig. 8. The MedFusion Interface displaying the input handling and real-time generated report

The test case has gone through the interface's end-to-end processing process. Fig. 8. In the left section of the interface are displayed the input files (the audio file and the radiology image) that were uploaded into the interface for review by the processing engine. The right section of the interface contains the final Clinical Report that was generated by the Rule-Based

Fusion Engine in real-time. The "Patient History" section of this report is structured, as evidenced by the matching of the symptoms that were obtained from the LLM module and includes a statement of "Patient history indicates abdominal discomfort with Radiological imaging results indicate a chest x-ray showing a solid nodule in the left lung field." This statement shows how the system was able to combine both the subjective complaints from the patient as well as the objective findings described in the radiologic image.

IX. RESULTS AND ANALYSIS

For testing the MedFusion framework, a mixed method was performed by applying both quantitative measures and case study analysis. Although traditional measures have been employed to evaluate each module of the system, the effectiveness of the system as a whole was tested via test cases to guarantee the consistency and correctness of the output reports. The evaluation findings point out the critical balance between performance measurements and clinical relevance from a qualitative perspective. The conventional measures for evaluating systems such as ROUGE and BLEU show acceptable levels of performance; however, a more thorough investigation shows that the use of transformers (ViT + GPT-2) leads to outputs that have more meaning than those generated using CNN models.

An audio file in MP3 format contains a conversation between the physician and patient, explains symptoms of an acute gastroenteritis. He reports that he had his last meal in a Chinese restaurant five days ago. The patient's medical history indicates he has asthma and is a university student, living in shared accommodation with other students. An X-ray image was taken as visual input that shows the posteroanterior (PA) view of a chest radiograph with an L marked on the left side. The image is a typical view of the thoracic cavity, including lungs, ribs, mediastinum and diaphragm. The image is diagnosed by experts as having a solid nodule in the patient's left lung field. Through the Speech Processing Module, the audio has been processed to harvest strict clinical information, not simply a synopsis.

Deterministic fusion is very important in guaranteeing that the entire system functions reliably. Since this process ensures strict template-based integration of outputs, the system avoids cross-modality issues and minimizes error propagation from module to module. Contrary to the generative method of fusion, this approach ensures that an error made by one modality will not affect the results. Furthermore, the nearly instantaneous time taken for fusion of outputs (<0.01 seconds) emphasizes the efficiency of rules-based fusion over neural networks.

X. CONCLUSION

Even with these advantages, however, the use of rules for cross-modal fusion results in limitations in the capability of the model to acquire complex relationships within cross-modal data, since this may not be well-defined in the template design. While this makes the system safe and explainable, it might result in limited flexibility for more complex cases. Through the creation of an end-to-end AI network of integrations between how they review images, the MedFusion Project has effectively integrated these traditionally isolated processes into one seamless workflow through a series of optimized, sequential pipelines on home-grown hardware.

One significant problem with the existing evaluation paradigm is the lack of a standard benchmark database that includes synchronized audio recordings, images, as well as clinical reports. Therefore, evaluation in our study can be seen as qualitative rather than quantitative benchmark evaluation of the system. The next direction for future research will include building up such databases and performing evaluation based on the opinion of clinicians.

Although the present use of MedFusion shows the potential for multiple modes of clinical AI, research has ways in which this use could be improved by using a more enhanced segmentation and area selection, enhancing the exact placement of clinical findings followed with the development of a system that allows AI to learn from the experience of experts.

REFERENCES

- [1] Z. Wang, M. Tang, L. Wang, X. Li, and L. Zhou, "A Medical Semantic-Assisted Transformer for Radiographic Report Generation," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2022*, 2022.
- [2] L. Wang, M. Ning, D. Lu, D. Wei, Y. Zheng, and J. Chen, "An Inclusive Task-Aware Framework for Radiology Report Generation," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2022*, 2022.
- [3] J. Li, S. Li, Y. Hu, and H. Tao, "A Self-guided Framework for Radiology Report Generation," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2022*, 2022.
- [4] Q. Chen and Y. Hong, "MedBLIP: Bootstrapping Language-Image Pretraining from 3D Medical Images and Texts," in *ACCV 2024*, 2025.
- [5] J. Wang, A. Bhalerao, and Y. He, "Cross-modal Prototype Driven Network for Radiology Report Generation," arXiv preprint arXiv:2207.04818, 2022.
- [6] Z. Zhang, B. Wang, W. Liang, et al., "Sam-guided enhanced fine-grained encoding with mixed semantic learning for medical image captioning," arXiv preprint arXiv:2311.01004, 2023.
- [7] R. Tanno, D. G. T. Barrett, A. Sellergren, et al., "Consensus, dissensus and synergy between clinicians and specialist foundation models in radiology report generation," arXiv preprint arXiv:2311.18260, 2023.
- [8] S. Singh, "Designing a Robust Radiology Report Generation System," arXiv preprint arXiv:2411.01153, 2024.
- [9] J. Rückert, L. Bloch, R. Brüngel, et al., "ROCOv2: Radiology Objects in Context Version 2, an Updated Multimodal Image Dataset," *Scientific Data*, 2024.
- [10] T. Tanida, P. Müller, G. Kaissis, and D. Rueckert, "Interactive and Explainable Region-guided Radiology Report Generation," in *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [11] M. V. Arisoy, A. Arisoy, and İ. Uysal, "A vision attention driven Language framework for medical report generation," *Scientific Reports*, 2025.
- [12] S. K. Nayak, M. Garanayak, S. K. Swain, S. K. Panda, and D. Godavathi, "An Intelligent Disease Prediction and Drug Recommendation Prototype by Using Multiple Approaches of Machine Learning Algorithms," *IEEE Access*, 2023.
- [13] B. M. Hassan and S. M. Elagamy, "Personalized medical recommendation system with machine learning," *Neural Computing and Applications*, 2025.
- [14] A. Kumar, G. K. Sharma, and U. M. Prakash, "Disease Prediction and Doctor Recommendation System using Machine Learning Approaches," *International Journal for Research in Applied Science and Engineering Technology*, 2021.
- [15] Q. Zhou, L. Su, L. Wu, and D. Jiang, "Deep Personalized Medical Recommendations Based on the Integration of Rating Features and Review Sentiment Analysis," *Wireless Communications and Mobile Computing*, 2021.
- [16] Krishnan, Sajitha, et al. "Improving accessibility in learning using visual question answering and eye gaze insights." Information and Communication Systems. CRC Press, 2026. 480-486.
- [17] Venugopalan, Manju, and Deepa Gupta. "An enhanced guided LDA model augmented with BERT based semantic strength for aspect term extraction in sentiment analysis." Knowledge-based systems 246 (2022): 108668.
- [18] Akhilesh, P., et al. "Automated Short Answer Grading using Machine Learning and Deep Learning Models." 2025 3rd World Conference on Communication & Computing (WCONF). IEEE, 2025.
- [19] Poojitha, D., et al. "Multi-Modal Detection and Classification of Diabetic Retinopathy." 2025 International Conference on Electronics, Computing, Communication and Control Technology (ICECCC). IEEE, 2025.
- [20] Narayani, Mudumala Varnika, et al. "Leveraging Machine Learning to Predict Disease Risk Levels from Multiple Eye Condition Indicators." 2025 5th International Conference on Pervasive Computing and Social Networking (ICPCSN). IEEE, 2025.